

First Linear Regression Model

Importing the basic libraries

```
In [59]: 1 import pandas as pd
          2 import numpy as np
          3 import seaborn as sns
          4 import matplotlib.pyplot as plt
          5 %matplotlib inline
          6 import warnings
          7 warnings.filterwarnings(action='ignore', category=FutureWarning)
```

Importing the dataset from sklearn The name of the dataset is boston

```
In [2]: 1 from sklearn.datasets import load_boston
```

```
In [60]: 1 df=load_boston()
```

Checking the keys of dataset

```
In [6]: 1 df.keys()
```

```
Out[6]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename', 'data_module'])
```

Description of the dataset

```
In [8]: 1 print(df.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset  
-----
```

```
**Data Set Characteristics:**
```

```
    :Number of Instances: 506
```

```
    :Number of Attributes: 13 numeric/categorical predictive. Median Value (a  
ttribute 14) is usually the target.
```

```
    :Attribute Information (in order):
```

```
        - CRIM      per capita crime rate by town  
        - ZN        proportion of residential land zoned for lots over 25,000  
sq.ft.  
        - INDUS     proportion of non-retail business acres per town  
        - CHAS      Charles River dummy variable (= 1 if tract bounds river; 0  
otherwise)  
        - NOX       nitric oxides concentration (parts per 10 million)  
        - RM        average number of rooms per dwelling  
        - AGE       proportion of owner-occupied units built prior to 1940  
        - DIS       weighted distances to five Boston employment centres  
        - RAD       index of accessibility to radial highways  
        - TAX       full-value property-tax rate per $10,000  
        - PTRATIO   pupil-teacher ratio by town  
        - B          $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of black peop  
le by town  
        - LSTAT     % lower status of the population  
        - MEDV      Median value of owner-occupied homes in $1000's
```

```
    :Missing Attribute Values: None
```

```
    :Creator: Harrison, D. and Rubinfeld, D.L.
```

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/> ([https://a
rchive.ics.uci.edu/ml/machine-learning-databases/housing/](https://archive.ics.uci.edu/ml/machine-learning-databases/housing/))

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

```
.. topic:: References
```

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [12]:

```
1 print(df.data)
```

```
[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

```
In [11]: 1 print(df.target)
```

```
[24.  21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15.  18.9 21.7 20.4
 18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
 18.4 21.  12.7 14.5 13.2 13.1 13.5 18.9 20.  21.  24.7 30.8 34.9 26.6
 25.3 24.7 21.2 19.3 20.  16.6 14.4 19.4 19.7 20.5 25.  23.4 18.9 35.4
 24.7 31.6 23.3 19.6 18.7 16.  22.2 25.  33.  23.5 19.4 22.  17.4 20.9
 24.2 21.7 22.8 23.4 24.1 21.4 20.  20.8 21.2 20.3 28.  23.9 24.8 22.9
 23.9 26.6 22.5 22.2 23.6 28.7 22.6 22.  22.9 25.  20.6 28.4 21.4 38.7
 43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
 18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22.  20.3 20.5 17.3 18.8 21.4
 15.7 16.2 18.  14.3 19.2 19.6 23.  18.4 15.6 18.1 17.4 17.1 13.3 17.8
 14.  14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
 17.  15.6 13.1 41.3 24.3 23.3 27.  50.  50.  50.  22.7 25.  50.  23.8
 23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
 37.9 32.5 26.4 29.6 50.  32.  29.8 34.9 37.  30.5 36.4 31.1 29.1 50.
 33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50.  22.6 24.4 22.5 24.4 20.
 21.7 19.3 22.4 28.1 23.7 25.  23.3 28.7 21.5 23.  26.7 21.7 27.5 30.1
 44.8 50.  37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29.  24.  25.1 31.5
 23.7 23.3 22.  20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
 29.6 42.8 21.9 20.9 44.  50.  36.  30.1 33.8 43.1 48.8 31.  36.5 22.8
 30.7 50.  43.5 20.7 21.1 25.2 24.4 35.2 32.4 32.  33.2 33.1 29.1 35.1
 45.4 35.4 46.  50.  32.2 22.  20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
 21.7 28.6 27.1 20.3 22.5 29.  24.8 22.  26.4 33.1 36.1 28.4 33.4 28.2
 22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21.  23.8 23.1
 20.4 18.5 25.  24.6 23.  22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
 19.5 18.5 20.6 19.  18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
 22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25.  19.9 20.8 16.8
 21.9 27.5 21.9 23.1 50.  50.  50.  50.  50.  13.8 13.8 15.  13.9 13.3
 13.1 10.2 10.4 10.9 11.3 12.3  8.8  7.2 10.5  7.4 10.2 11.5 15.1 23.2
  9.7 13.8 12.7 13.1 12.5  8.5  5.  6.3  5.6  7.2 12.1  8.3  8.5  5.
 11.9 27.9 17.2 27.5 15.  17.2 17.9 16.3  7.  7.2  7.5 10.4  8.8  8.4
 16.7 14.2 20.8 13.4 11.7  8.3 10.2 10.9 11.  9.5 14.5 14.1 16.1 14.3
 11.7 13.4  9.6  8.7  8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
 14.1 13.  13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20.  16.4 17.7
 19.5 20.2 21.4 19.9 19.  19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
 16.7 12.  14.6 21.4 23.  23.7 25.  21.8 20.6 21.2 19.1 20.6 15.2  7.
  8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
 22.  11.9]
```

Getting the name of the features

```
In [13]: 1 print(df.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

Creating the dataset from array of boston data frame

```
In [18]: 1 data=pd.DataFrame(df.data,columns=df.feature_names)
```

Adding the price columns to our dataset

```
In [20]: 1 data['Price']=df.target
```

Viewing the top 5 rows of dataset

```
In [21]: 1 data.head()
```

Out[21]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

Checking the dypes of the columns

```
In [24]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    float64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  Price       506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

Checking the statistics of the numerical columns

In [25]: 1 data.describe()

Out[25]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.00
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.79
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.10
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.12
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.10
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.20
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.18
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.12



In [34]: 1 data.shape

Out[34]: (506, 14)

EDA

Checking the null values in the dataset

In [29]: 1 data.isna().sum()

Out[29]:

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
Price	0

dtype: int64

Checking the correlation of the dataset

In [37]:

```
1 data.corr()
```


Out[37]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670	0.
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408	-0.
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027	0.
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176	-0.
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230	0
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246	-0.
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881	0.
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000	-0.
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.494588	1.
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.534432	0.
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.232471	0.
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.291512	-0.
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.496996	0.
Price	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.249929	-0.



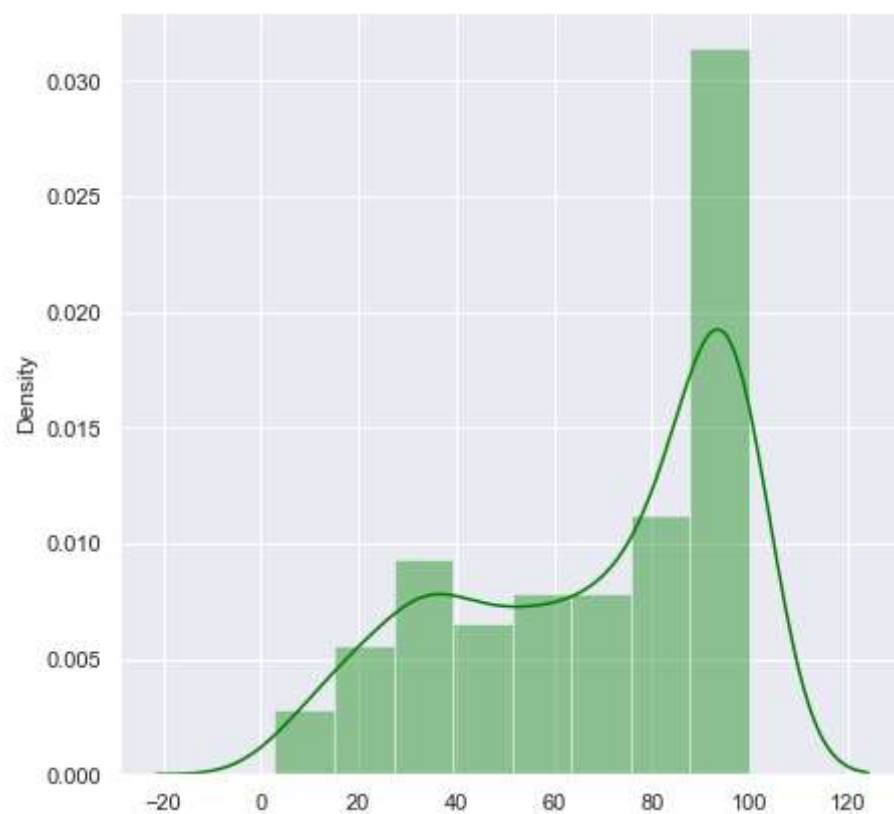
```
In [42]: 1 sns.set(rc={'figure.figsize':(10,10)})
2         sns.heatmap(data.corr(),annot=True,cmap="YlGnBu",linewidths=.5,cbar=False)
```

Out[42]: <AxesSubplot:>





```
In [61]: 1 sns.set(rc={'figure.figsize':(7,7)})  
        2 sns.distplot(x=data["AGE"],hist=True,kde=True,color="GREEN")
```

Out[61]: <AxesSubplot:ylabel='Density'>




```
In [57]: 1 sns.scatterplot(data=data,x="CRIM",y="Price")
```

```
Out[57]: <AxesSubplot:xlabel='CRIM', ylabel='Price'>
```




```
In [67]: 1 sns.regplot(data=data,x="DIS",y="Price",scatter=True,marker="+",color="green"
```

```
Out[67]: <AxesSubplot:xlabel='DIS', ylabel='Price'>
```




```
In [72]: 1 sns.lmplot(data=data,x="RM",y="Price")
```

```
Out[72]: <seaborn.axisgrid.FacetGrid at 0x1b3bf0d0d00>
```




```
In [82]: 1 sns.jointplot(data=data,x="CRIM",y="Price",kind="reg",color="green",height=1
```

```
Out[82]: <seaborn.axisgrid.JointGrid at 0x1b3c05a5700>
```




```
In [84]: 1 sns.set(rc={'figure.figsize':(7,5)})
        2 sns.boxplot(data['TAX'],color="green",linewidth=2,palette="Set3")
```

Out[84]: <AxesSubplot:xlabel='TAX'>



```
In [90]: 1 sns.violinplot(x=data["RM"],palette="Set1")
```

Out[90]: <AxesSubplot:xlabel='RM'>



Splitting the features and target in x and y variables

```
In [92]: 1 x=data.iloc[:, :-1]
        2 y=data.iloc[:, -1]
```

In [93]:

```
1 x
```

Out[93]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.

506 rows × 13 columns



In [96]:

```
1 y
2
```

Out[96]:

```
0    24.0
1    21.6
2    34.7
3    33.4
4    36.2
...
501   22.4
502   20.6
503   23.9
504   22.0
505   11.9
Name: Price, Length: 506, dtype: float64
```

Splitting the dataset into train and test part

In [97]:

```
1 from sklearn.model_selection import train_test_split
```

In [125]:

```
1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_sta
```

```
In [152]: 1 x_train
```

```
Out[152]: array([[ -0.45098379,  0.59163708, -0.84258861, ...,  0.55738858,
                  0.37807205, -0.82757883],
                 [ 7.93012362, -0.50215615,  1.06054133, ...,  0.80335661,
                 -3.79668193, -0.32645016],
                 [ 0.46202474, -0.50215615,  1.06054133, ...,  0.80335661,
                  0.40378237, -0.6918886 ],
                 ...,
                 [-0.45261131,  0.04474047, -0.44053957, ..., -1.65632372,
                  0.44498343,  1.06745852],
                 [-0.45759249, -0.50215615,  0.44279445, ..., -1.06600044,
                  0.41098987, -0.26631472],
                 [-0.42648556, -0.50215615, -0.51243885, ...,  0.50819497,
                  0.44498343, -0.93705617]])
```

```
In [101]: 1 x_train.shape
```

```
Out[101]: (354, 13)
```

```
In [102]: 1 x_test.shape
```

```
Out[102]: (152, 13)
```

```
In [105]: 1 y_test.shape
```

```
Out[105]: (152,)
```

```
In [106]: 1 y_train.shape
```

```
Out[106]: (354,)
```

Standardize or feature scaling the dataset

```
In [107]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [122]: 1 scaler=StandardScaler()
```

```
In [126]: 1 x_train=scaler.fit_transform(x_train)
```

```
In [127]: 1 x_test=scaler.transform(x_test)
```



```
In [128]: 1 x_train
```

```
Out[128]: array([[ -0.45098379,  0.59163708, -0.84258861, ...,  0.55738858,
                   0.37807205, -0.82757883],
                  [ 7.93012362, -0.50215615,  1.06054133, ...,  0.80335661,
                  -3.79668193, -0.32645016],
                  [ 0.46202474, -0.50215615,  1.06054133, ...,  0.80335661,
                   0.40378237, -0.6918886 ],
                  ...,
                  [-0.45261131,  0.04474047, -0.44053957, ..., -1.65632372,
                   0.44498343,  1.06745852],
                  [-0.45759249, -0.50215615,  0.44279445, ..., -1.06600044,
                   0.41098987, -0.26631472],
                  [-0.42648556, -0.50215615, -0.51243885, ...,  0.50819497,
                   0.44498343, -0.93705617]])
```

```
In [129]: 1 x_test
```

```
Out[129]: array([[ 0.37877319, -0.50215615,  1.06054133, ...,  0.80335661,
                   0.42916997, -0.11520515],
                  [-0.46870399,  2.99798219, -1.37229555, ..., -2.9353575 ,
                  -0.01317667, -0.55774031],
                  [-0.43463958, -0.50215615, -0.58140346, ..., -0.32809634,
                   0.44498343,  2.865354  ],
                  ...,
                  [-0.44178915,  0.37287844, -1.10670841, ..., -1.80390454,
                   0.44498343, -1.13750764],
                  [-0.4670962 ,  1.24791302, -0.65477008, ..., -0.47567716,
                   0.44498343, -1.34104298],
                  [-0.31388118, -0.50215615, -0.4009216 , ...,  1.19690546,
                   0.41959583,  1.18001973]])
```

Model Building

```
In [130]: 1 from sklearn.linear_model import LinearRegression
```

```
In [131]: 1 lr=LinearRegression()
```

```
In [132]: 1 lr
```

```
Out[132]: LinearRegression()
```

```
In [134]: 1 lr_model=lr.fit(x_train,y_train)
```

```
In [135]: 1 lr_model.score(x_train,y_train)
```

```
Out[135]: 0.7431215456774967
```

```
In [157]: 1 Linear_regression_coefficient=lr_model.coef_
```

```
In [158]: 1 Linear_regression_coefficient.transpose
```

```
Out[158]: <function ndarray.transpose>
```

```
In [159]: 1 Linear_regression_coefficient
```

```
Out[159]: array([-0.6208519 ,  0.89604528, -0.4181019 ,  0.85794528, -1.98345156,
                2.34054146, -0.14708338, -2.8644969 ,  2.15413705, -1.58410776,
                -1.74439973,  0.6305477 , -3.22010917])
```

```
In [181]: 1 Lr_coefficient=pd.DataFrame(data=df.feature_names,columns=["Independent facto
2
```

```
In [182]: 1 Lr_coefficient["coefficient"]=Linear_regression_coefficient
```

```
In [183]: 1 Lr_coefficient
```

```
Out[183]:
```

	Independent factors	coefficient
0	CRIM	-0.620852
1	ZN	0.896045
2	INDUS	-0.418102
3	CHAS	0.857945
4	NOX	-1.983452
5	RM	2.340541
6	AGE	-0.147083
7	DIS	-2.864497
8	RAD	2.154137
9	TAX	-1.584108
10	PTRATIO	-1.744400
11	B	0.630548
12	LSTAT	-3.220109

```
In [137]: 1 lr_model.intercept_
```

```
Out[137]: 22.331355932203394
```

```
In [184]: 1 predicted_values=lr_model.predict(x_test)
```

```
In [185]: 1 predicted_values
```

```
Out[185]: array([21.90897572, 32.36829283,  9.38919345, 16.40673353, 17.80964232,
 31.83838312, 25.10363218, 15.4942598 , 21.82825591, -3.63190569,
 26.12960431, 15.57300292,  5.61225053,  5.58756072, 25.41154332,
 34.70503462, 26.17912943, 19.13532445, 23.91967422, 14.91252997,
 39.53465438, 11.07641307, 36.58914352, 26.00446715, 38.64469005,
 25.17973575, 21.75528189, 18.96547913, 18.27571802, 18.60093947,
 24.62357132, 23.66620392, 29.6987949 , 24.08585329,  0.50581275,
 24.63764742, 25.21913509, 12.19902726, 39.4812705 , 32.23454473,
 23.75474746,  7.056712 , 20.39810217, 21.0026853 , 31.32729178,
  7.46193071, 12.70824342, 31.32832609, 22.40993904, 35.817382 ,
 12.81513925, 20.71658302, 18.48252207,  7.65314991,  6.48378445,
 40.45412148, 24.95009747, 24.17943728, 23.04271387,  7.56345617,
 22.86100568,  9.73479018, 32.957889 , 14.06778493, 28.52717573,
 17.20171167,  3.61911076, 28.62629983, 19.42447388, 18.72979294,
 19.36051351, 27.88976576, 21.11155756, 27.95264386, 34.21722447,
 20.00321067, 13.39987071, 24.70136312, 16.70346939, 22.70991552,
 18.90523529, 17.48644847, 18.59401509, 10.09315724, 16.78988769,
 10.9323577 , 17.072616 , 20.4084587 , 20.34209532, 19.17179969,
 27.40674633,  7.73791552, 20.11982554,  5.18969717, 20.2172659 ,
  5.18968724, 17.55397823, 27.00736521, 23.00701384, 20.51409442,
 24.44825342, 16.23006422, 24.80902733,  5.95769186, 31.26003941,
 21.69010225, 30.42222709, 31.82079134, 21.90982483, 17.40866927,
 29.98143286, 39.96512278, 27.59995934, 22.03687566, 22.24622018,
 14.94470564, 21.28339441, 19.3278602 , 41.8844113 , 21.35566086,
 22.53418169, 28.98276109, 25.20250497, 16.47346987, 41.69742145,
 18.17821943, 13.59812528, 24.6032745 , 16.28450795, 28.54125121,
 13.22639708, 26.69967006, 30.09184714, 23.06307131, 33.64391574,
 34.87452978, 19.21096706, 20.22625076, 13.38442001, 19.58013994,
 12.97469066, 35.1171442 , 16.93050452, 25.13534002, 27.26167899,
 22.67914379, 14.24375693, 23.22069174, 14.87303988, 34.55202866,
 36.1305225 , 13.89590962])
```

```
In [187]: 1 Comparsion_table=pd.DataFrame(data=predicted_values,columns=['Predicted valu
```

```
In [189]: 1 Comparsion_table["Actual Value"]=y_test
```

```
In [191]: 1 Comparsion_table["Difference"]=Comparsion_table["Actual Value"]-Comparsion_t
```

```
In [200]: 1 Comparsion_table.dropna().head()
```

```
Out[200]:
```

	Predicted values	Actual Value	Difference
0	21.908976	24.0	2.091024
3	16.406734	33.4	16.993266
4	17.809642	36.2	18.390358
6	25.103632	22.9	-2.203632
9	-3.631906	18.9	22.531906

Performance Metrics


```
In [210]: 1 from sklearn.metrics import mean_squared_error
2 from sklearn.metrics import mean_absolute_error
3 print("Mean squared error:",mean_squared_error(y_test,predicted_values))
4 print("Mean absolute error:",mean_absolute_error(y_test,predicted_values))
5 print("Root mean squared error:",np.sqrt(mean_squared_error(y_test,predicted
```

Mean squared error: 31.829631155557482
Mean absolute error: 3.9079661456255192
Root mean squared error: 5.641775532184658

Assumptions Of Linear Regression

```
In [211]: 1 sns.regplot(x=y_test,y=predicted_values)
```


Out[211]: <AxesSubplot:xlabel='Price'>



```
In [212]: 1 residuals=y_test-predicted_values
```


```
In [216]: 1 sns.displot(residuals,kind="kde",palette="Set1",color="green")
```

```
Out[216]: <seaborn.axisgrid.FacetGrid at 0x1b3c6af2ac0>
```



```
In [220]: 1 sns.scatterplot(x=predicted_values,y=residuals,color='green',)
```

```
Out[220]: <AxesSubplot:ylabel='Price'>
```



R square and adjusted R square

```
In [221]: 1 from sklearn.metrics import r2_score
2 R_score=r2_score(y_test,predicted_values)
3 print("R score value is",R_score)
```

R score value is 0.7215519718844172

```
In [223]: 1 Adjusted_r_score=1 - (1-R_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]
2 print("Adjusted R score value is",Adjusted_r_score)
```

Adjusted R score value is 0.6953213605401957

Ridge regression model

```
In [225]: 1 from sklearn.linear_model import Ridge  
          2 ridge=Ridge()
```

```
In [229]: 1 ridge_model=ridge.fit(X_train,y_train)
```

```
In [228]: 1 ridge_predict=ridge.predict(X_test)
```

```
In [230]: 1 ridge_model.coef_
```


```
Out[230]: array([-8.96773369e-02,  3.87293612e-02, -9.46326932e-02,  3.27853539e+00,  
                -8.83310616e+00,  3.70498171e+00, -1.32155805e-02, -1.23740801e+00,  
                2.25516417e-01, -1.00845466e-02, -7.84860767e-01,  7.35041802e-03,  
                -5.07459245e-01])
```

```
In [231]: 1 ridge_model.intercept_
```

```
Out[231]: 30.29949040952099
```


```
In [232]: 1 sns.regplot(x=y_test,y=ridge_predict)
```

```
Out[232]: <AxesSubplot:xlabel='Price'>
```




```
In [233]: 1 ridge_residuals=y_test-ridge_predict
          2 sns.displot(ridge_residuals,kind="kde",palette="Set1",color="green")
```

Out[233]: <seaborn.axisgrid.FacetGrid at 0x1b3c7eed6d0>




```
In [236]: 1 sns.scatterplot(x=ridge_predict,y=ridge_residuals,color='green',)
```

```
Out[236]: <AxesSubplot:ylabel='Price'>
```



Performance Metrics

```
In [237]: 1 print("Mean squared error:",mean_squared_error(y_test,ridge_predict))
2 print("Mean absolute error:",mean_absolute_error(y_test,ridge_predict))
3 print("Root mean squared error:",np.sqrt(mean_squared_error(y_test,ridge_pre
```

```
Mean squared error: 32.0345368225147
Mean absolute error: 3.9014189572909572
Root mean squared error: 5.659906078948192
```

Lasso regression model

```
In [238]: 1 from sklearn.linear_model import Lasso
```

```
In [239]: 1 lasso_model=Lasso()
```

```
In [242]: 1 lasso_model=lasso_model.fit(x_train,y_train)
```

```
In [243]: 1 lasso_predict=lasso_model.predict(x_test)
```

```
In [244]: 1 lasso_predict
```

```
Out[244]: array([21.34943756, 30.45318044,  9.76791574, 20.28915417, 17.71113013,
 27.45637077, 26.66319632, 17.10112721, 22.57985801,  1.7296673 ,
 20.58479405, 17.87865192,  4.34340966,  6.70983829, 22.71114893,
 34.28455434, 23.055947  , 19.39217948, 24.49889928, 15.66732815,
 35.42870375, 13.60752019, 32.90886105, 24.98379635, 36.14242032,
 29.03925819, 21.93532581, 21.71817967, 20.24747986, 20.50831643,
 24.86664646, 24.07003804, 28.8757406 , 25.6785744 , -0.23400937,
 24.88197726, 25.55953235, 11.9363952 , 36.60086295, 26.74081754,
 25.1929427 , 10.52873993, 19.15795595, 21.47613761, 30.72956898,
  7.79566921, 13.75848843, 29.15598704, 20.78588524, 32.55445146,
 14.25879118, 20.78376397, 17.16959343, 13.37098929, 15.63037829,
 37.22159371, 25.0004319 , 18.0836259 , 22.62264165, 10.20743311,
 23.25511343, 13.11813923, 29.77533454, 13.91016337, 25.49115206,
 15.98976512,  4.57956761, 29.78260633, 18.52792707, 19.12915633,
 21.2179911 , 25.61629936, 21.83311757, 29.12291699, 32.11381314,
 19.40982872, 16.12273523, 25.9968222 , 18.11940996, 26.45392971,
 18.66219764, 13.49719305, 20.24261533, 13.4594494 , 18.5513846 ,
 10.93173044, 16.79905557, 22.29027312, 20.12925555, 18.65810201,
 23.43081403, 10.24298825, 20.55058669,  8.65794007, 19.43321998,
  6.13141912, 18.65069388, 26.13931311, 22.35657609, 19.83791901,
 24.14245176, 18.01044292, 24.11018372,  6.37749415, 29.36638489,
 24.39914647, 27.89469674, 31.00315572, 22.67316443, 17.63904305,
 34.00548455, 35.37768265, 26.62176649, 22.7781386 , 24.97210281,
 19.50897172, 21.62116718, 22.70604851, 35.89213029, 22.14704538,
 22.93953403, 26.57059325, 28.8755948 , 17.78095571, 38.17535806,
 17.83236195, 21.2435862 , 21.80465299, 17.91024311, 26.2268157 ,
 15.20023204, 25.82043688, 27.49505856, 22.09921939, 30.15073758,
 32.54023852, 19.47952826, 20.87293377, 13.42168405, 21.76504344,
 11.8022726 , 34.9518832 , 20.38607318, 25.20425326, 26.84212068,
 19.36409275, 13.79327497, 18.89432554, 17.60815573, 30.11257116,
 29.4974732 , 15.27235538])
```

```
In [245]: 1 lasso_model.coef_
```

```
Out[245]: array([-0.          ,  0.          , -0.          ,  0.09154745, -0.          ,
 2.43878411, -0.          , -0.          , -0.          , -0.11045338,
-1.07955341,  0.          , -3.33388011])
```

```
In [246]: 1 lasso_model.intercept_
2
```

```
Out[246]: 22.331355932203394
```

```
In [256]: 1 Lasso_model_coefficient=pd.DataFrame(data=df.feature_names,columns=["Independ
```

```
In [257]: 1 Lasso_model_coefficient["lasso coefficient"]=lasso_model.coef_
```


```
In [258]: Lasso_model_coefficient
```

```
Out[258]:
```

	Independent factors	lasso coefficient
0	CRIM	-0.000000
1	ZN	0.000000
2	INDUS	-0.000000
3	CHAS	0.091547
4	NOX	-0.000000
5	RM	2.438784
6	AGE	-0.000000
7	DIS	-0.000000
8	RAD	-0.000000
9	TAX	-0.110453
10	PTRATIO	-1.079553
11	B	0.000000
12	LSTAT	-3.333880

```
In [259]: 1 sns.regplot(x=y_test,y=lasso_predict)
```

```
Out[259]: <AxesSubplot:xlabel='Price'>
```



```
1 ### Performance Metrics
```

```
In [260]: 1 print("Mean squared error:",mean_squared_error(y_test,lasso_predict))
2 print("Mean absolute error:",mean_absolute_error(y_test,lasso_predict))
3 print("Root mean squared error:",np.sqrt(mean_squared_error(y_test,lasso_pre
```

Mean squared error: 40.3546502317852
Mean absolute error: 4.28976184314201
Root mean squared error: 6.35253100990347


Elastic Net regression model

```
In [267]: 1 from sklearn.linear_model import ElasticNet
```

```
In [269]: 1 en=ElasticNet()
2 Elasticnet_model=en.fit(x_train,y_train)
3 Elasticnet_predict=en.predict(x_test)
```

```
In [272]: 1 sns.regplot(x=y_test,y=Elasticnet_predict)
```

Out[272]: <AxesSubplot:xlabel='Price'>



Performance Metrics

```
In [273]: 1 print("Mean squared error:",mean_squared_error(y_test,Elasticnet_predict))
2 print("Mean absolute error:",mean_absolute_error(y_test,Elasticnet_predict))
3 print("Root mean squared error:",np.sqrt(mean_squared_error(y_test,Elasticne
```

Mean squared error: 43.90854807474513
Mean absolute error: 4.295235872781503
Root mean squared error: 6.626352546819791

Comparsion of coeffiecent between linear,Ridge,Lasso regression

```
In [264]: 1 Comparsion_of_coeffiecent=pd.DataFrame(data=df.feature_names,columns=["Indep
```

```
In [270]: 1 Comparsion_of_coeffiecent["Linear regreesion coefficient"]=Linear_regression  
2 Comparsion_of_coeffiecent["Ridge regression coefficient"]=ridge_model.coef_  
3 Comparsion_of_coeffiecent["Lasso regression coefficient"]=lasso_model.coef_  
4 Comparsion_of_coeffiecent["ElasticNet regression coefficient"]=Elasticnet_mo
```

```
In [271]: 1 Comparsion_of_coeffiecent
```

Out[271]:

	Independent factors	Linear regreesion coefficient	Ridge regression coefficient	Lasso regression coefficient	ElasticNet regression coefficient
0	CRIM	-0.620852	-0.089677	-0.000000	-0.117033
1	ZN	0.896045	0.038729	0.000000	0.062132
2	INDUS	-0.418102	-0.094633	-0.000000	-0.467268
3	CHAS	0.857945	3.278535	0.091547	0.522033
4	NOX	-1.983452	-8.833106	-0.000000	-0.320978
5	RM	2.340541	3.704982	2.438784	2.105497
6	AGE	-0.147083	-0.013216	-0.000000	-0.017915
7	DIS	-2.864497	-1.237408	-0.000000	-0.000000
8	RAD	2.154137	0.225516	-0.000000	-0.000000
9	TAX	-1.584108	-0.010085	-0.110453	-0.369758
10	PTRATIO	-1.744400	-0.784861	-1.079553	-1.059870
11	B	0.630548	0.007350	0.000000	0.365535
12	LSTAT	-3.220109	-0.507459	-3.333880	-2.083438

From this table we can conclude that how overfitting and underfitting can be reduce with ridge and lasso regreesion. In ridge regression in cost function we take lamba and square of slope in plus of linear regreesion but In Lasso regreesion in cost function we take lamba and absolute of slope in plus of linear regreesion In Elastic net regression in cost function we take addition of linear,ridge and lasso cost function

```
In [ ]: 1
```